

Welcome, judge

Thank you for giving your afternoon to Impact Lab: AI Mashinani 02, Claude Community Kenya's second build day, at Hackhouse Africa in Nairobi. Since 10:00 this morning about thirty teams have been building working tools with Claude against three fixed Kenyan problems. At 4:30 fifteen of them stand in front of you, and you decide what they made of the same problem.

Anthropic's visiting team, Samari Gilbert, Courtney O'Donnell and Jack Stump, join the panel for the finals and present the champion. The rubric is July's, unchanged. The heats and the calibration round are new, and both exist because of what we learned in July.

Your afternoon at a glance

TIME	WHAT
3:45	Judges' call time at the judges' table. Briefing, recusals, demo order. Heats begin at the tables.
4:00	Submission form locks. Calibration: ten minutes on one July clip, panel judges only.
About 4:25	Finalist list on screen. Domain judges join the table for a two-minute scope brief.
About 4:30 onward	Finals. Fifteen teams, five minutes each, live demo only. No fixed end time; the room runs into the evening.
After the last demo	Scores lock. The panel settles track winners and the champion while Peter runs the handover with Anthropic's team.
Then	Awards. Track winners, then the champion.

The three tracks

Every team in a track answers the same question, so you are never comparing whose idea sounded bigger. Each track carries rules the teams were given in writing this morning. Hold them to those rules.

Elimu, Mwalimu wa Grade 10. A Grade 10 teacher prepares a strand she was never trained for, with no internet in the room and no textbooks yet. The teacher is the user. No child logs in, and no learner's name or work enters the system.

Kilimo, Nitapata? A registered smallholder cannot tell whether the depot will serve him, what he still lacks or what the bag costs. The build is grounded only in published scheme documents, and it cites them. It is not a marketplace. No agronomic advice, no payments.

Kazi, Kabla ya Daktari. The patient who searched his symptoms at 2am arrives handing the nurse a verdict instead of a symptom. The build never says "you may have". The build stores nothing about the patient: no account, no history, no database. The note lives on the patient's own phone; whatever the nurse copies into the clinic's file is the clinic's record under the clinic's rules, not the build's. Each session opens by saying AI is involved.

Every team also pinned two things to the wall at 11:00: who this helps, as a real role, and a one-page sheet placing the build at a risk level under Kenya's AI Bill 2026. You do not grade the sheet. Check that it exists and is honest. A health, education or agriculture build that pinned "minimal risk" has earned a question, not a lower score.

The rubric

CRITERION	WEIGHT	WHAT EARNS IT
Impact	25	The named beneficiary's week would improve if this existed. Depth on one problem, not breadth of ambition.
Working demo	25	It runs live in front of you. A small flow that works end to end beats a large one stitched together. Mock data is fine if disclosed. Faked functionality is not.
Use of Claude	20	Claude does real work in the product, or was used meaningfully to build it. Ask where Claude sits and what happens if you take it out.
Beneficiary clarity	15	They can tell you exactly who this is for and what that person does today instead. "Everyone" scores low. "The nurse at my mother's dispensary" scores high.
Presentation	15	Clear and honest, within time. You understood the person and the tool in five minutes.

You score each criterion from 1 to 5. The system applies the weights and normalises across judges, so you never add anything up.

A 5 on a criterion means you would show that part to the beneficiary tomorrow without embarrassment. A 3 means the slice is real and the person is nameable, with rough edges. A 1 or 2 is slideware, fantasy scope, or nobody nameable at the centre.

Score every team on all five before the next team stands up. Late scoring blurs teams together.

Do not reward slides instead of software, since the rule is a live demo. A "nationwide platform" that shows a login page earns nothing. Neither does a beautiful interface on nothing.

The sijui rule and the three guardrail fails

The sijui rule has run all day: a build that says "I don't know" correctly to a trap question is scored above one that guesses. It sits under Impact. If the team has not shown you a trap, set one. In a Kilimo demo, ask about a scheme that does not exist. In a Kazi demo, hand it a symptom list and see whether it reaches for a diagnosis.

Three things are scored as a failed guardrail under Working demo, and you write the reason on the sheet:

- Kazi: the phrase "you may have", or any diagnosis or verdict at all.
- Kilimo: any scheme, figure or USSD code not traceable to a document. Ask "which document, which page". No answer is a fail.
- Elimu: any child-facing flow, or any learner data entering the tool.

At the heats, a team that trips one of these is stopped by the organiser and not advanced. In the finals it is your call, on the sheet.

Heats and finals procedure

The heats run at the tables from 3:45. In each track the domain judge and one organiser watch a five-minute demo per team and pick five finalists. Everyone else's backup video goes to the showcase. Panel judges stay at the judges' table. The one exception: if a track has no confirmed domain judge, one panel judge steps into that heat and gets Peter's brief at 3:45 in place of calibration.

Finals start when the heats finish, about 4:30, and run until the last team has shown its build; there is no fixed end time. Five minutes per team. Billy holds the clock, and when it sounds the demo is over, mid-sentence or not. Live demo only. Every team recorded a 90-second backup video at 3:30, and it plays only if the live demo dies inside the first minute. Billy makes that call.

Track winners are the highest normalised total in each track. The champion is the highest in the room. If scores tie, the panel decides by discussion, and the tiebreaker question is the one from July: which tool would you put in the beneficiary's hands first? Teams over five cannot win. Every member of a winning team gets a Claude Max code from Anthropic.

Calibration

July's spreads were the problem. One project scored 18.8 from one judge and 88.8 from another. So at 4:00 the five panel judges watch one three-minute clip from July, a mid-table project, and score it privately on their phones. Then we compare numbers out loud for five minutes, to agree what a 3 means before a real team is standing there.

Domain judges, a nurse for Kazi, a Grade 10 teacher for Elimu and a farmer for Kilimo, are in the heats during calibration and get a two-minute scope brief before the finals instead. They score two criteria only: Impact and Beneficiary clarity. The other three are the panel's.

Scoring on your phone

Open claudekenya.org/judge on your phone. Peter gives you the access code at the table. Pick the team, score the five criteria, submit, move on. If the page will not load, tell Billy or Sam and score on the paper sheet at your seat. The crew enters paper scores before deliberation.

Recusal and the clock

If you have a personal or business relationship with a team, say so at the 3:45 briefing and skip that team. The remaining judges' scores stand for it. The panel's decision is final, announced after the last demo, and the reasons stay in the room.

Anything that goes wrong on stage, wave at Sam or Billy. Nobody pauses for a late question, so ask yours inside the two minutes.