

CLAUDE COMMUNITY KENYA

EVENT REPORT

Claude Conversations, Nairobi

**Saturday 29 August 2026,
Blockchain Centre NBO**

Prepared by Peter Kibet, Claude Code Ambassador, Kenya

PUBLISHED 2 SEPTEMBER 2026

REVISED 4 SEPTEMBER 2026

Summary

About 15 people spent Saturday afternoon at Blockchain Centre working the three questions from the Claude Conversations kit as one table, in turns. The plan was 50 people at eight tables. With one table there were no cards on a wall and no vote.

Four things kept coming back. Frontier AI is priced for another market, and the figures the room put beside the plan price are in the asks section. The access gap for Kenyan children is opening, and one attendee has already built an offline tutor with a teacher in the loop. Kenyan data is missing from the models, and nobody could say who would own it once captured. The first rung of work is going, from junior developer hiring to the nurse whose patients arrive with an app's verdict.

Those four became three fixed problems for Wednesday's Impact Lab and three asks for Anthropic's Policy team.

Event at a glance

ITEM	DETAIL
Date	Saturday 29 August 2026
Venue	Blockchain Centre NBO, Argwings Kodhek Road, Nairobi
Host	Peter Kibet, Claude Code Ambassador for Kenya
Format as planned	7 to 8 tables of 6 to 8, one prompt per table, read-outs, cards on the wall, dot vote, 2:00 to 5:00 PM
Format as run	One table, all three prompts in turn, round-robin, roughly 3:45 to 5:30 PM
Attendance	About 15, against a plan for 50
Prompts from the kit	1. What does AI mean for my job? 2. What is AI doing to me, my kids, my community? 3. Who decides the rules?
Feeds	Impact Lab: AI Mashinani 02, Wednesday 2 September 2026, Hackhouse Africa

Who was in the room

A nurse. A farmer from Sultan Hamud now building a coffee-traceability platform. A specialist working in a One Health context. A builder of small language models for edge devices, who brought a working Raspberry Pi prototype. Staff from two banks. Developers, a data-validation analyst and a UX designer. A former senior software engineer now in DevRel. Someone who runs two communities for young people with disabilities. A student. The host uses Claude daily and runs his own farm-management product, MkulimaOS.

What the room said

What does AI mean for my job

“if you’re still talking to Anthropic, at least tell them to adjust for purchasing power.”

The nurse opened. Patients now describe their symptoms to an app overnight and arrive at the clinic with a conclusion instead of a question. An app can tell someone “it has told me I have cancer”, and she cannot dismiss that, but “this person is just coming with what they’ve read online all night.” She asked for a way of narrowing the information between the app and the consultation.

The DevRel speaker said that three years ago “I knew that software engineers were going to be the first attack.” Hiring for software developers “has dropped drastically” and companies now want AI-native builders. The UX designer read the same moment as opportunity: “It tells me to pivot fast, super, super fast.” The displacement cases the room agreed on were Chegg, whose traffic collapsed once students could ask an AI, and prompt engineering, a job title that was over as fast as it began “because it was a service.”

The longest thread of the afternoon was about cost. The room argued about the cause, whether Claude is expensive or everything in Kenya is, and never about the fact that frontier AI is priced for another market. The ask came from the room: “if you’re still talking to Anthropic, at least tell them to adjust for purchasing power.”

AI, me, kids, community

A participant who builds farming software asked the farmer what AI could do farm to market. The farmer wanted an agentic system that walks a farmer “from the idea of what to farm, how to farm it, what to purchase, all the way to selling.” Fancy dashboards are close to useless for farmers who are not formally educated. Something they can talk to and direct would work.

On children the room agreed without dissent that the gap is opening: “there is also a 9 year old somewhere who is still counting sticks,” set against a nine-year-old elsewhere building an agent, both expected to compete in the same world for the same resources. One speaker put the cause in policy: “it’s more like an infrastructural problem.” The edge-models builder described the prototype: a Raspberry Pi serving its own Wi-Fi with no backhaul, a small model over CBC curriculum content and government scheme information, and a teacher curating what students see. “It doesn’t need to be powerful, right? Because the scope is narrow.”

On community, one speaker said African traditions are oral, so there is little for models to train on, and “we’re losing whatever makes us us” to imported popular culture. She then asked: when we solve our own community’s problems, “will we just train the AI so that those solutions will come back and be sold to us as another subscription?”

Several people said heavy Claude use had made them antisocial. One had joined the ambassador programme “out of loneliness of using Claude in the house.”

Who decides the rules

Three answers, no winner. The first was top-down: “government, because governments make policy, policy affects corporations, corporations make the products,” with company reputation as the only check. The second disagreed outright: “I think the person who decides the rules is the consumer.” Whatever earns money gets built, and if a government banned something, someone would build it anyway. The third split it three ways: government first, then the corporation, then the user, whose rule is what they choose to put in.

The round ended on the instructions a user never sees. Every model carries “a whole book of instructions ahead of your own instruction.” The last word was half a joke: “Mr. Dario will make the rules.”

Findings

No vote was held, so the order below carries no weight.

1

Frontier AI is priced for a market this room is not in.

The room argued about the cause, whether Claude is expensive or everything is, and never about the fact. The figures sit with the ask in the asks section. This goes to Anthropic’s Policy team.

2 The access gap for children is opening, and the working design keeps an adult in the loop.

The Raspberry Pi tutor keeps a teacher between the model and the child: “I wanted to have a human in the loop, which is the teacher.” A one-day team can build the offline software layer with the teacher as the user. Devices and power stay out of scope. This became Track 1.

3 Kenyan data is missing from the models, and nobody said who would own it.

The One Health specialist said Claude is more precise in her domain but that African data is thin because most of it is on paper: “there’s information on our end, but it’s not out there, it’s not accessible.” A recording interface is easy. Consent and revocable ownership are the hard part. This goes to the Policy team with the pricing ask.

4 The first rung of work is going, and the nurse’s waiting room shows where it lands.

Junior hiring has dropped drastically and nobody disputed it. The nurse’s patients hand her a verdict instead of a symptom, so her first job is to reopen the question. The buildable piece is at her end: something that narrows what the patient brings to the consultation without diagnosing. This became Track 3.

How the tracks were decided

The kit’s closing step was not run; the time was gone. On the morning of Monday 31 August the host cut three tracks from the findings, one problem per track, so that judges compare execution. Two shapes were dropped before the tracks reached the team: a parent-facing homework explainer, because it would carry a child’s identifiable homework into the system, and the farm-to-market agent, because the agent must not act on the farmer’s behalf.

Later that day the volunteer team reviewed the tracks. The central objection: “you’re saying we are giving them a solution to build rather than a problem.” Track 1 as then written specified the exact output, so every team would build the same thing. The counter-position, also in the room: “we need to create problems that will be easy to judge.” The team stated a goal for each track apart from its mechanism. Elimu: “to improve the quality of education children get through the plan.” Kilimo: “the goal is to give fair value to the farmer.” Kazi: “the problem is self-diagnosis.”

The host rewrote the tracks on the team’s side. The problem and its rules are fixed and binding. The build is open. Under each track sits one worked answer so builders can see the scale of thing that fits in a day, marked “one answer, not the answer.” The final wording goes to the approved builders before doors open on 2 September, and the blast says in the host’s own words that the team “pushed back hard on one thing” and that “they were right.”

Impact Lab tracks

Everything below except the build is binding.

TRACK 1 **Elimu** Mwalimu wa Grade 10

PROBLEM

A Grade 10 teacher prepares a strand she was never trained for. There is no internet in the room and the textbooks have not arrived.

GOAL

The quality of education the children in that class get, not the software.

FIXED RULES

The teacher is the user, no child logs in, and no learner's name or work enters the system. Judges ask one question: would a real Grade 10 teacher use this for tomorrow's lesson, offline?

ONE ANSWER, NOT THE ANSWER

She picks a strand and a lesson length and gets a 40-minute plan, five questions with a marking scheme and one no-materials activity, as an offline pack. One attendee's Raspberry Pi tutor is the offline lane.

TRACK 2 **Kilimo** Nitapata?

PROBLEM

A registered smallholder cannot tell whether the depot will serve him, what he still lacks, what the bag costs or whether a loan fits, and he would rather ask a phone than travel to an officer.

GOAL

Fair value to the farmer, and a wasted trip prevented rather than caused.

FIXED RULES

Grounded only in published scheme documents, it cites the document, and it names the missing step. No agronomic advice, no dosage, no payments, no autonomous action. One warning: most agri-tech already tries to cut out the middleman, so a produce marketplace will not win here.

ONE ANSWER, NOT THE ANSWER

A USSD, SMS, voice or chat line, in Kiswahili or English.

TRACK 3 **Kazi** Kabla ya Daktari

PROBLEM

The patient has already diagnosed himself. He searched his symptoms at 2am, argued with a chatbot until it agreed with him, and arrives handing the nurse a verdict instead of a symptom.

GOAL

Stop self-diagnosis turning into misdiagnosis, without replacing the clinician.

FIXED RULES

Never “you may have”, the clinic stores nothing, and each session opens by saying AI is involved. Judges ask whether the clinician got a better starting point than they would have had. A team may swap the nurse for another professional (a pharmacist, a vet, a paralegal) only if that professional is on the team and in the room.

ONE ANSWER, NOT THE ANSWER

A number the patient messages in the waiting room, four questions (what you searched, what it said, what you fear it is, what you actually feel), and a one-page note on the patient’s own phone to show the nurse.

Asks carried to Anthropic

The asks were put to Anthropic’s visiting team in person at the Impact Lab on 2 September. This report carries them in writing.

Pricing for purchasing power. A Max plan with tax costs about KES 27,000 a month. The gazetted minimum wage for a general labourer in Nairobi is reported as KES 18,047 under Legal Notice No. 109 of 2026, effective 1 May 2026 (press reports; gazette not independently checked). Hardware has moved the same way: 32 GB of DDR5 at KES 30,000, about four times last year’s price, and a 1 TB SSD at KES 35,000 to 40,000, roughly double what older stock sold for. Everything is expensive here. The ask, in the room’s words, is to “adjust for purchasing power.”

Data terms. Who owns Kenyan data used for training? The room did not get as far as terms. The host’s proposal, which nobody voted on: a licence that the contributor can revoke, and no cloning of identifiable voices, which the room raised as a live fraud problem aimed at parents.

A small model that runs offline. The edge-models builder’s Raspberry Pi tutor serves its own Wi-Fi with no backhaul: a small model over a curriculum knowledge base, with a teacher in the loop. On 29 August it ran on the Claude API while his free tokens lasted. What he wanted instead was a narrow model that performs well offline on the Pi: “It doesn’t need to be powerful, right? Because

the scope is narrow.” Four days later, at the Impact Lab, his team had an open-weight model from another lab running live on the Pi with no connection. The room asks for a small Claude-class model that can run offline on cheap hardware in a classroom or a clinic with no connectivity.

AI governance

The third prompt asked who decides the rules. Kenya's Senate is part-way through answering, so this page sets the room's answers beside the Bill and beside what the Impact Lab did with it.

WHAT THE ROOM SAID

Three answers, none of which won. Government, because policy shapes corporations and corporations make the products. The consumer, because whatever earns money gets built. Or a three-way split: government first, then the corporation, then the user. The last word was half a joke: "Mr. Dario will make the rules."

WHAT THE BILL SAYS

The Artificial Intelligence Bill, 2026 (Senate Bills No. 4 of 2026) had its first reading on 2 April 2026 and is not yet law. Clause 25 sorts AI systems into four tiers:

- Unacceptable risk: systems that pose severe threats. Prohibited outright.
- High risk: systems used in the critical sectors of healthcare, education, agriculture, finance, security, employment and public administration.
- Limited risk: moderate risk.
- Minimal risk: negligible risk.

A high-risk system owes three things: a risk assessment and a human rights impact assessment before deployment, with human oversight; transparency, traceability and explainability of its decisions, with records kept for five years; and registration on the Commissioner's public register, with an annual compliance report. Clause 28 binds every tier: disclose the system's nature, purpose and limitations, disclose how much of a decision is automated, and disclose what was done about bias.

WHAT THE IMPACT LAB DID WITH IT

Every team filled a one-page risk-level sheet by 11:00 and pinned it beside its "This helps" card. Health, education and agriculture are on the Bill's list of critical sectors, so nearly every build in the room landed at high risk, the expected tier for these tracks. The sheet asked each team to name its tier with a reason, then write down the three duties that would follow.

WHAT THIS ROOM ASKED OF ANTHROPIC'S POLICY TEAM

Price the plans for Kenyan purchasing power. Say who would own Kenyan data used for training. The host's opening proposal is a licence the contributor can revoke and no cloning of identifiable voices. And consider a small Claude-class model that runs offline on cheap hardware. The figures and the evidence are in the asks section.

Corrections welcome. Send them to Peter Kibet.